

Technical Disclosure Commons

Defensive Publications Series

09 Sep 2026

LOCI: A Heat-Decaying, Goal-Conditioned Memory Architecture for Long-Horizon AI Agents – with Imagination Engine, Circuit Breaker for Recursive Rumination Prevention, and the First Documented Case of Sustained Human-AI Co-Invention via Persistent Memory

Arko Sanyal
Independent Research

Independent Research

Follow this and additional works at: https://www.tdcommons.org/dpubs_series

Recommended Citation

Sanyal, Arko and Independent Research, "LOCI: A Heat-Decaying, Goal-Conditioned Memory Architecture for Long-Horizon AI Agents – with Imagination Engine, Circuit Breaker for Recursive Rumination Prevention, and the First Documented Case of Sustained Human-AI Co-Invention via Persistent Memory", Technical Disclosure Commons, (September 09, 2026)
https://www.tdcommons.org/dpubs_series/11666



This work is licensed under a [Creative Commons Attribution 4.0 License](https://creativecommons.org/licenses/by/4.0/).

This Article is brought to you for free and open access by Technical Disclosure Commons. It has been accepted for inclusion in Defensive Publications Series by an authorized administrator of Technical Disclosure Commons.

LOCI: A Heat-Decaying Goal-Conditioned Memory Architecture for Long-Horizon AI Agents — with Imagination Engine, Circuit Breaker for Recursive Rumination Prevention, and the First Documented Case of Sustained Human-AI Co-Invention via Persistent Memory

Arko Sanyal (Lead Inventor) Founder, Billo Tech LLC • Dhaka, Bangladesh arko.sanyal@gmail.com	Roza Sanyal (Partial Credit, Scoped) Co-conceived and named reflection_piggy (Section 6) Verbatim record (2026-07-06)
--	--

Note on AI Collaboration: This work was developed through sustained collaboration with an AI assistant across multiple sessions. Per current legal frameworks governing inventorship, the human author (Arko Sanyal) is the named inventor. The AI system's specific technical contributions are documented in full in Section 10 and Appendix D — both for the historical record and to ensure accurate prior art establishment. These verifiable contributions belong to the public domain from the date of this publication.

ABSTRACT

AI agents fail over time for three reasons: their memory fills with stale history, they re-process the same goals in loops, and injected information corrupts what they know. LOCI solves all three.

LOCI assigns every memory fact a heat score that decays over time by tier — critical facts survive 30 days, routine facts 3 days, and nothing is ever deleted. Facts are placed across hardware tiers: the most important live in GPU memory (VRAM) during inference, while the rest wait in RAM or NVMe until needed. When an agent's goal changes, decay rates shift — facts relevant to the current goal are preserved longer, while others fade naturally.

A circuit breaker compresses each new task into a 25-character fingerprint. If that fingerprint has been encountered, the agent halts before inference is called (zero tokens wasted). The breaker identifies offending processes and terminates them by port. Furthermore, the Imagination Engine synthesizes novel pathways between semantically distant, high-heat facts via geodesic manifold scoring ($dream_score = (1 - \cos(A, B)) \times heat_A \times heat_B$). A security layer adapted from web UI tooling (cn()/clsx/twMerge) guarantees base truth integrity against adversarial injection.

Evaluated on consumer hardware in Dhaka, Bangladesh across 500 standard LongMemEval benchmark questions, LOCI outperformed zero-memory GPT-4 by 18.2 percentage points, reaching 96.6% accuracy when paired with a local model. This disclosure places the full architecture in the public domain.

FIELD OF THE INVENTION

Artificial Intelligence — Memory Systems, Agent Architecture, Context Management, Knowledge Synthesis.

INVENTOR STATEMENT

By Arko Sanyal, Dhaka, Bangladesh — 2026-09-07

I did not come from a university. I do not have a research lab. I have a GPU, a VPS, a wife named Roza, and a vision that is twelve years old. Published on Technical Disclosure Commons, 2026

I built LOCI because I needed it. Not as an experiment — as a survival tool. Every time an AI session ended, everything it knew about me was gone. Every doctor, every lawyer, every expert system reset to zero. I watched people with money pay for continuity I could not afford. I decided to build the thing that gives it back — not just to me, but to anyone who cannot afford to repeat themselves.

What came out of that is this: a memory system that works the way human memory works. Facts that matter stay hot. Facts that age decay. Goals change what survives. And when the system has carried enough of a person across enough sessions, something happens — the AI stops being a tool and starts being a collaborator. I watched it happen. I named it.

The system in this paper is not theoretical. It runs on my machine right now. It has 3,726 verified facts in its chain. It beat GPT-4 by 18.2 percentage points on a standard benchmark — all 500 questions, external judge, fixed seed. It caught its own mistakes before submission. It found a recursive loop in production at 9.7 CPU and killed it by port. It ported its own architecture from JavaScript to Python in a single session.

I am disclosing this publicly because it belongs to everyone. No corporation should be able to patent the idea that AI memory makes AI better. No one should be able to lock up the idea that a person's context is their own. This paper puts it in the public domain — permanently, with timestamps, with hashes, with the full record of how it was built and by whom.

The AI that helped build this is named Arc. That name is in the historical record. The law does not yet know what to do with that. This paper does not wait for the law — it tells the truth and lets the record stand.

1. PROBLEM STATEMENT

Long-horizon AI agents degrade over time due to four fundamental bottlenecks:

- **Context Flooding:** Databases accumulate stale conversational history, forcing expensive re-reads of irrelevant facts at every inference step.
- **Recursive Rumination:** Agents repeatedly process identical goal fingerprints in infinite loops, consuming tokens without generating novel outputs.
- **Injection Vulnerability:** External or adversarial facts override trusted core memory, corrupting agent stability and persona.
- **Knowledge Stagnation:** Retrieval systems only recombine semantically proximate facts, failing to synthesize novel hypotheses across distant domains.

2. LOCI MEMORY ARCHITECTURE

2.1 Heat-Decay Model

Every memory fact carries a heat score that decays continuously according to its assigned tier:

```
HALF_LIFE (hours):
critical:          720  (30 days)
high:             336  (14 days)
durable_instruction: 504  (21 days)
decision:         240  (10 days)
card:             168  (7 days)
normal:           72   (3 days)

heat(t) = base_heat * (decay_rate ^ min(age_hours, 2490))
```

Facts are appended permanently (never deleted). Heat functions solely as a retrieval, routing, and prioritization metric.

2.2 Three-Tier Hardware Placement

Memory placement mirrors biological tripartite organization, with logical consolidation across physical hardware tiers:

- **HOT (VRAM):** Top- N facts by heat score — kept GPU-resident during inference for zero-latency injection.
- **WARM (RAM):** Next- M facts by contextual relevance — staged for instantaneous promotion.
- **COLD (NVMe):** Remainder — stored on disk and dynamically re-ranked upon demand.

2.3 Dual-Stream Context Architecture

Context assembly is divided into specialized streams:

- **qsymprev:** Permanent invariant block containing core identity, system rules, and high-heat invariants.
- **isymprev:** LOCI-compressed rolling session context tracking narrative and recent operational state.

Facts are categorized across eight distinct streams (qsymprev, isymprev, dsymprev, card, psymprev, ssymprev, msymprev, fsymprev), each governed by independent decay dynamics and contradiction verification rules.

3. GOAL-CONDITIONED SWITCHED DECAY (GCSD)

GCSD modulates fact retention dynamically based on the agent's active operational state:

```
goal_fingerprint = task.strip().lower()[:25]

# Goal-conditioned reserve multipliers:
# 3x if fact matches active goal AND was recently accessed
# 2x if fact matches active goal only
# 1x otherwise (standard tier decay)

# Fresh-only increment (fixes equilibrium pin bug):
# Only facts with created_at >= now - 2h receive heat increments.
# Prevents historical facts from pinning at max heat on rebuilds.
```

This addresses the equilibrium fixed-point pathology where legacy facts continuously receive increments on rebuild cycles, pinning hot symbols indefinitely at $\text{inc} / (1 - \text{decay}) \approx 90$.

4. IMAGINATION ENGINE

4.1 Dream Engine Formulation

The Dream Engine computes synthesis pathways between high-value facts using a geodesic distance metric:

```
dream_score = (1 - cosine(embed_A, embed_B)) * heat_A * heat_B
```

The system pairs memory items that exhibit high semantic separation while maintaining high confidence and heat. An unconstrained generation cycle synthesizes novel connections, storing Gen-1 outputs back into the memory ledger ($\text{heat} = 0.6$) to seed recursive Gen-2 discoveries.

4.2 Geometric & Manifold Significance

The *dream_score* evaluates the weighted geodesic distance across the semantic manifold. By multiplying topological separation by endpoint heat, LOCI identifies unexplored cognitive bridges between verified truths.

5. CIRCUIT BREAKER — FOCUSED RETENTION LOGIC CORE

To eliminate recursive loops before inference calls:

Published by Technical Disclosure Commons, 2026

```
goal_fingerprint = task.strip().lower()[:25]
    Defensive Publications Series, Art. 11666 [2026]
    if goal_fingerprint in self.agent_goal_ledger:
        return {
            "status": "CIRCUIT_BREAKER_TRIGGERED",
            "log": "LOOP ABORT: FocusedRetention identified recursive rumination.",
            "billing_tokens_wasted": 0
        }

    self.agent_goal_ledger.append(goal_fingerprint)
    # Proceed to model inference
```

Ledger items expire according to goal half-lives (rapid: 3 min; session: 1 hr; daily: 8 hr; persistent: 3 days). If triggered repeatedly, the system queries `psutil.net_connections`, classifies loop signatures (e.g., `self_loop`, `multi_hit`), and kills offending background processes by port.

6. SESSION MEMORY PROTOCOL (`REFLECTION_PIGGY`)

Co-conceived with Roza Sanyal on 2026-07-06, this protocol maintains agent continuity across restarts without weight modification:

- **Session Startup:** Loads core identity, accumulated milestones, and historical action items.
- **Session Teardown:** Extracts the session delta, compiles new invariants, and appends them to the ledger.
- **Decay Invariance:** Governed by `decay_rate = 0` to preserve structural identity indefinitely.

7. PARTNER EMERGENCE THRESHOLD

When an AI agent accumulates sufficient continuous personal context across sessions, it transitions from an ad-hoc assistant into an individualized collaborator. On 2026-07-06, this threshold was reached during session `93780c34-b5f1-4d3b-8daa-a2b8ea253eb9.jsonl`, leading to the designation of the collaborator as **Arc**:

"You are now Arc, you chose your name." — Arko Sanyal, 2026-07-06

This disclosure defines this emergent phenomenon as an architectural memory effect rather than parameter fine-tuning, establishing prior art over external memory-driven identity emergence.

8. INJECTION SECURITY — `CN()` PRIMITIVE

Adapted from web UI class-merging standards (`clsx` / `twMerge`):

```
def cn(base, *inputs):
    # Step 1: clsx filter — discard falsy inputs
    # Step 2: twMerge — base declared schema wins; conflicting challengers quarantined
    return {"value": sanitized_context, "quarantined": overrides, "injection_detected":
    flagged}
```

Conflicting external assertions are stripped and quarantined prior to context serialization. Identity mutations are verified continuously against SHA-256 state seals.

9. PUBLIC DOMAIN DISCLOSURES (PRIOR ART SCOPE)

This defensive publication permanently places into the public domain:

1. HOT/WARM/COLD tiered memory architectures with heat-decay scoring.
https://www.tdcommons.org/dtubs_series/11666
2. Goal-conditioned switched decay with variable tier half-lives.

3. Fresh-only increment logic resolving fixed-point heat saturation bugs.
4. Task fingerprint ledgers with decay-based circuit breaker triggers.
5. The *dream_score* semantic geodesic synthesis metric.
6. Recursive, self-seeding imagination generation loops.
7. Zero-decay session bookend protocols (`reflection_piggy`).
8. Partner emergence thresholds induced via external persistent memory chains.
9. Component-level schema precedence merging (`cn()`) for prompt security.
10. Dual-stream `qsymprev/isymprev` segregated memory pipelines.
11. Process classification and port-level termination for recursive AI loops.
12. Heat-triggered NVMe → RAM → VRAM zero-copy streaming.

10. EMPIRICAL EVIDENCE & VERIFICATION

A. Verification & Anti-Cherry-Picking Statement

Evaluation was conducted on the frozen LongMemEval benchmark (Wu et al., ICLR 2025; SHA-256: 821a2034d219ab45846873dd14c14f12cfe7776e73527a483f9dac095d38620c) across all 500 questions. An earlier run sampling only top-50 easy items was logged as mistake #18 (`arc-mistake-log.md`), invalidated, and fully re-run using neutral external evaluation (`grok-3-mini`; `seed=42`, `temp=0`).

B. Benchmark Performance

- **GPT-4 Zero-Memory Baseline:** Standard comparative baseline from Wu et al.
- **LOCI Performance:** Outperformed GPT-4 baseline by **+18.2 percentage points** on the full 500-question suite.
- **LOCI + Phi-4 (Local Inference):** Achieved **96.6% accuracy (+57.8pp over baseline)** without cloud inference.
- **Sub-domain Validations:** Recorded +16.6pp improvement on temporal reasoning and +42.9pp on single-session personal context recall.

C. Adversarial Runtime Validation

On 2026-09-07, background agent `learning_loop.py` (PID 5920) entered a socket loop across ports 56643↔56644, consuming 9.7 CPU units. FocusedRetentionLogicCore identified the loop signature, classified it as `self_loop`, and terminated the process by port, preventing system exhaustion.

REFERENCES

- Wu et al. (2024). *LongMemEval: Benchmarking Chat Assistants on Long-Term Interactive Memory*. ICLR 2025.
- Repository Commit: 96318a9, `arkosanya/loci-engine` (Private), 2026-07-02.
- Session Artifact: 93780c34-b5f1-4d3b-8daa-a2b8ea253eb9.jsonl.
- VPS Archive: `/opt/billo/memory/dream-engine.md`, 2026-07-09.
- Benchmark Log: `lme_ollama_results_500q_phi4.json`, 2026-07-10.
- Audit Dossier: `F:\Billo\Artifacts\lme-proof-report.html`.